

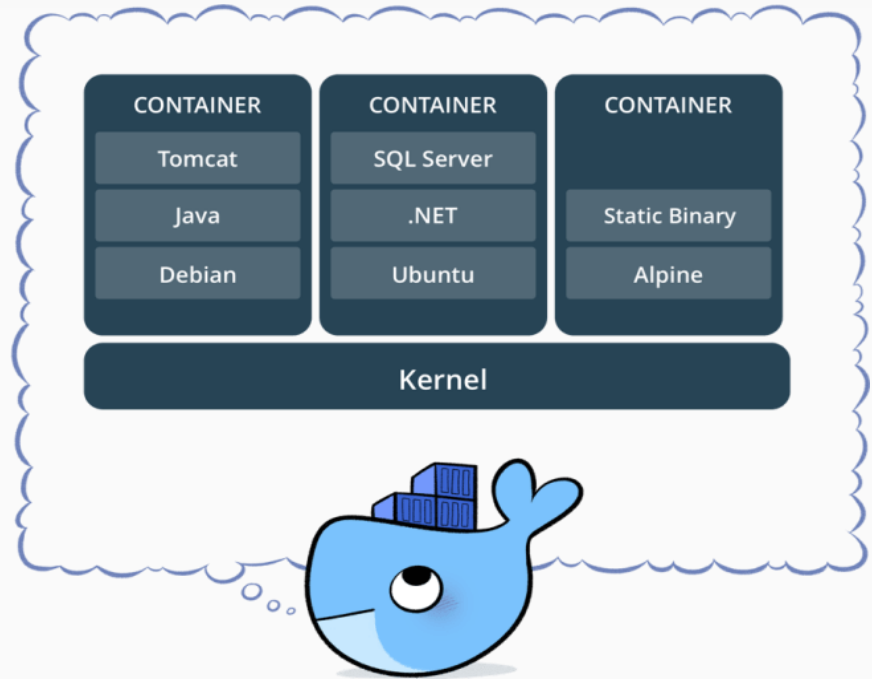
# Docker and Containers

An overview

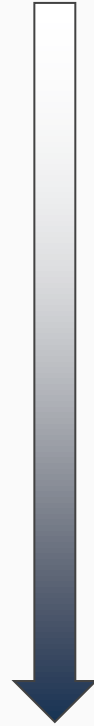


# What is Docker?

“ Docker is an open platform for developers and system administrators to build, ship and run containerized applications.”



# Major Infrastructure Shifts



Mainframe to PC (90's)

---

Bare Metal to Virtual (00's)













---

Datacenter to Cloud (10's)

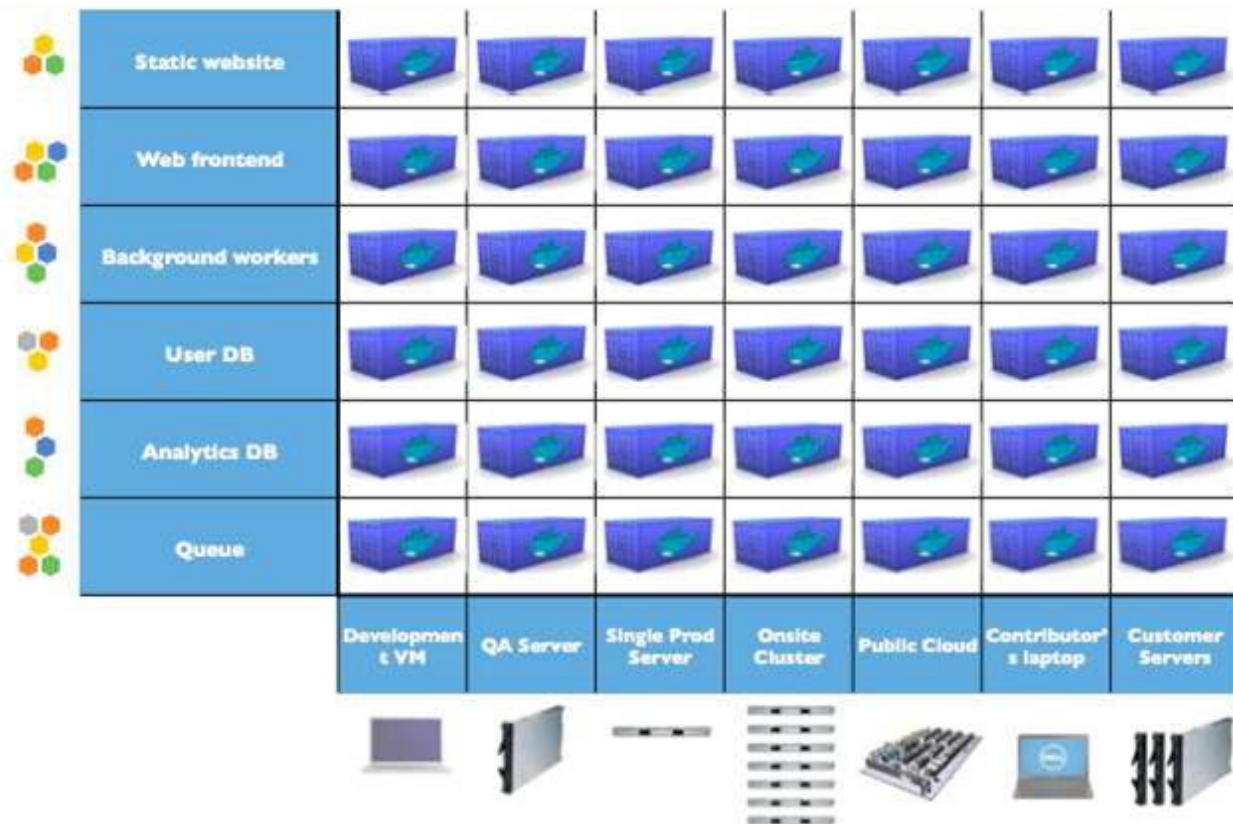
---

Host to Container  
(Serverless)

# Matrix from Hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

# Solving the matrix: Container



# Benefits of Docker:

## Speed

- Develop Faster
- Build Faster
- Test Faster
- Deploy Faster
- Update Faster
- Recover Faster

## Portability

- Less dependencies between process layers
- Ability to move between infrastructures

## Efficiency

- Less OS Overhead
- Improved VM Density

# Docker Editions



## Docker Community Edition

- Free; community-supported product for delivering a container solution
- Intended for: Software development and test



## Docker Enterprise Edition

- Subscription based commercially supported products for delivering a secure software supply chain
- Certified on specific platforms
- Intended for: Production Deployments and Enterprise Customers

# Docker on Linux: Setup

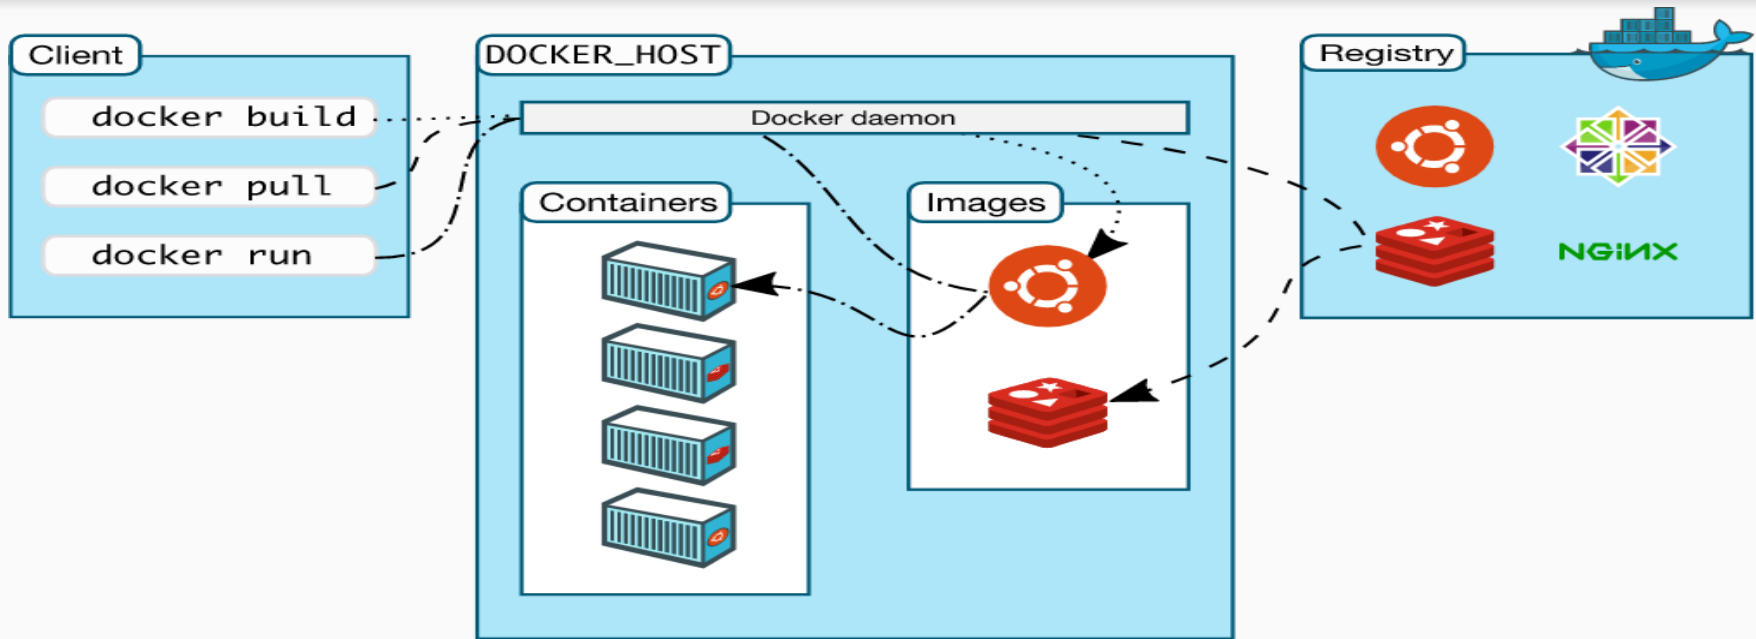
- Best native experience
- Three ways to install:
  - Script, store or ***docker-machine***
- get.docker.com script (latest Edge release)
  - ***curl -sSL https://get.docker.com/ | sh***
- [store.docker.com](https://store.docker.com) has instructions for each distro
- RHEL officially only supports Docker EE (paid), but CentOS will work
- Installing in a VM, Cloud Instance, all are the same process
- May not work for unlisted distros (Amazon Linux, Linode Linux, etc.)
- Don't use pre-installed setups (Digital Ocean, Linode, etc.)



# Installing Docker on Ubuntu

- [Installation Steps](#)
- [Register with Docker Hub](#)

# Docker Architecture:



Source: <https://docs.docker.com/engine/docker-overview/>

# Docker Architecture: (Contd.)

- Client-Server architecture having 3 main components:
  1. Docker Client
  2. Docker Host
  3. Docker Registry
- Docker Objects:
  - Images
  - Containers
  - Volumes
  - Networks

# Docker Architecture: (Contd.)

## **Docker Client:**

- The interface through which users interact with the docker using the docker commands.
- 2 basic ways to interact with Docker:
  - Docker CLI
  - Docker APIs
- Docker Client can communicate with more than one daemon

# Docker Architecture: (Contd.)

## **Docker Host:**

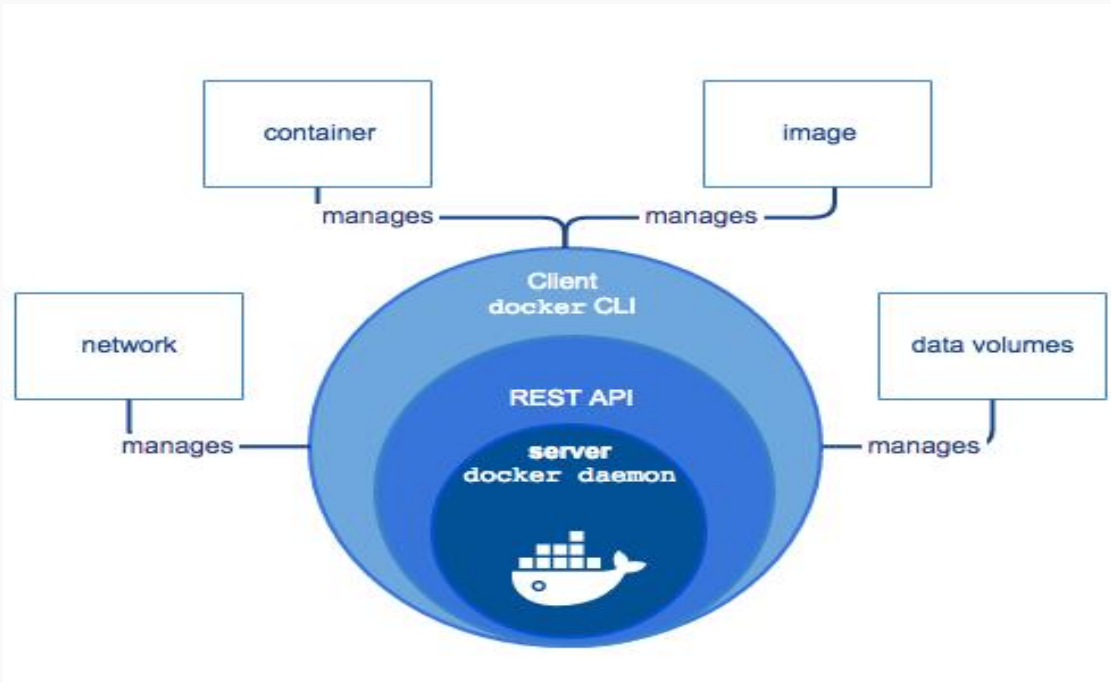
- Runs docker daemon which listens to and performs actions requested by Docker Client.
- Manages Docker objects such as images, containers, networks and volumes
- The Docker client and daemon can run on same systems or can connect docker client to a remote docker daemon
- Docker daemon is responsible for building dockerfiles, generating images and running containers

# Docker Architecture: (Contd.)

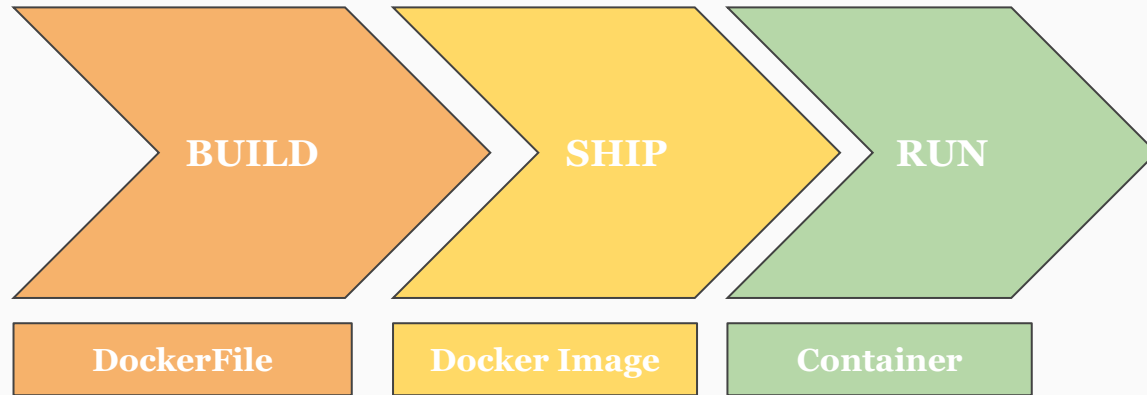
## Registry:

- Simplest component of docker architecture
- A universal place to store docker images and make them available to others
- Docker Hub is a public registry that anyone can use (<https://hub.docker.com/>)
- Docker is configured, by default, to look for docker images in Docker Hub
- *docker pull* and *docker push* are the commands used to pull and push the images from and to the configured registry

# Docker Components

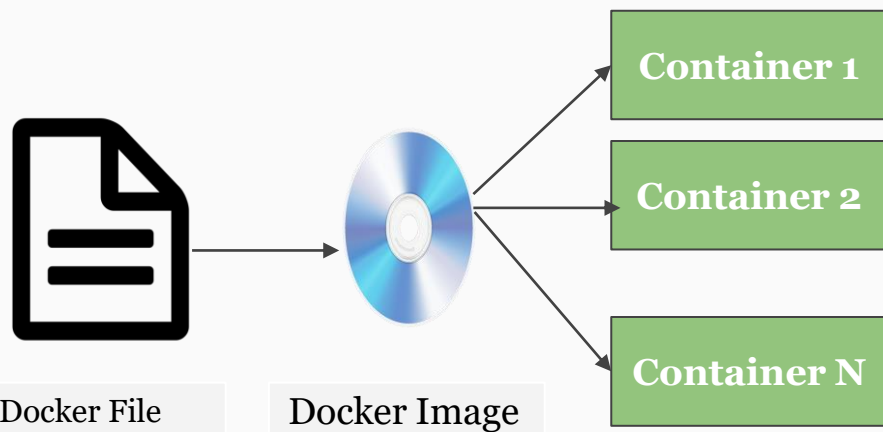


# Containerization Stages:





# Containerization Stages (Contd.):



- Multiple containers can be instantiated from same Docker image.
- One Dockerfile corresponds to one Docker Image

# Image vs Container

- An image is an application we want to run
- A container is an instance of that image running as a process
- Images: building blocks of the container
- You can have many containers running off the same image

# Let's Try:

```
docker container run -p 80:80 nginx
```

1. Download image 'nginx' from Docker Hub
2. Started a new container from that image
3. Opened port 80 on the host IP
4. Routes that traffic to the container IP, port 80

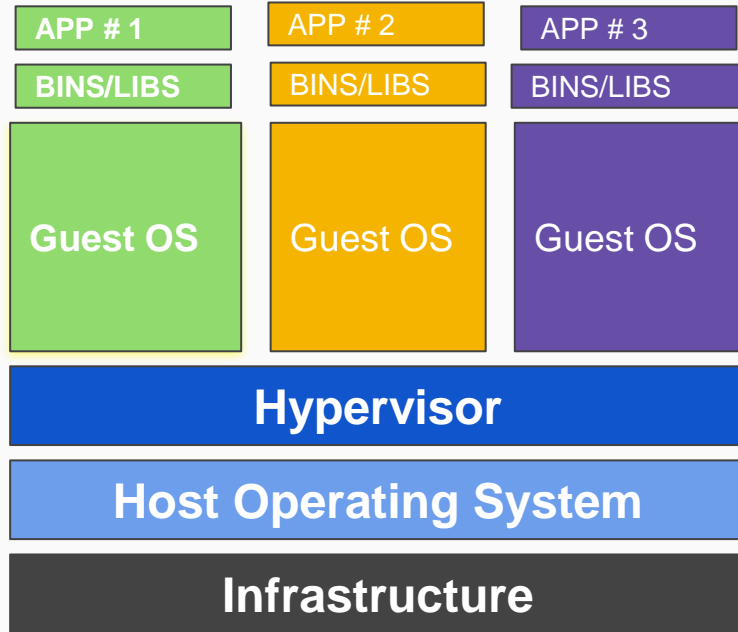
# How 'docker container run' works?

- Looks for specified image locally in image cache.
- If it does not find the image, then it looks in remote image repository (Docker Hub by default)
- Downloads the latest version of image (nginx: latest by default)
- Creates new container based on that image and prepares to start
- Gives it a virtual IP on a private network inside docker engine
- Opens up port 80 on host and forwards to port 80 in container
- Starts container by using CMD in the image Dockerfile

# Virtual Machines vs Container

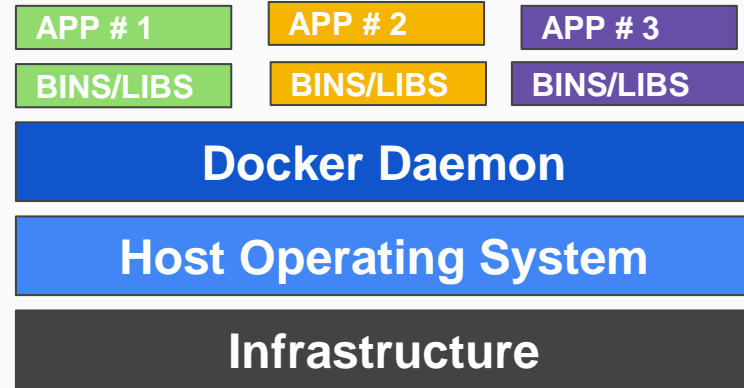
# Virtual Machines: vs Container

## Infrastructure Level Construct



**Virtual Machines**

## Application Level Construct



**Docker Containers**

## Virtual Machine vs Container

- Virtual machines are an abstraction of physical hardware turning one server into many server.
- Multiple VMs can run on single machine through Hypervisor
- Each VM has full copy of the OS
- Takes minutes to start
- Think of containers as isolated processes
- Containers are abstraction at the application layer that packages the code and dependencies together
- Each container runs in isolation with other containers;
- multiple containers can run on same machine sharing OS kernel with other containers
- Container can be up and running in milliseconds

## Virtual Machine vs Container (Contd...)

Some readings on why docker is faster than virtual machine

<https://www.backblaze.com/blog/vm-vs-containers/>



# Assignment 1: Manage Multiple Containers

- Run a **nginx**, a **mysql**, and a **httpd** (apache) server
- Run all of them **--detach (or -d)**, name them with **--name**
- Nginx should listen on **80:80**, httpd on **8080:80**, mysql on **3306:3306**
- When running mysql, use the **--env option (or -e)** to pass in **MYSQL\_RANDOM\_ROOT\_PASSWORD=yes**
- Use ***docker container logs*** on mysql to find the random password it created on startup
- Clean it up all with ***docker container stop*** and ***docker container rm*** (both the commands can accept multiple names and IDs)
- Use ***docker container ls*** to ensure everything is correct before and after clean up.
- Refer [docs.docker.com](https://docs.docker.com) and ***--help*** for any help

# What's happening within the container?

- To get process list of one container
  - ***docker container top***
- To get the configuration details of one container
  - ***docker container inspect***
- To get performance stats for all containers
  - ***docker container stats***
- Start new container interactively
  - ***docker container run -it***
- Run additional commands in existing container
  - ***docker container exec -it***
- Different Linux distros in container

# Docker Networks

# Section Overview:

- Review of ***docker container run -p***
- Concept of “batteries included but removable”
  - Defaults work but you can change a lot under the hood
- Quick port docker container port <container>
- Learn concepts of Docker Networking
- Understand how network packets move around Docker

# Docker Network Drivers: Basics

- When you start container, we are connecting to a particular Docker Network in the background.
- The communication is managed by objects called network drivers.
- A docker network driver is a piece of software which handles container networking.
- Created using ***docker network*** command.
- **'bridge'** is the default network driver. Usually used when your application runs in standalone containers that need to communicate.

# Docker Network: Defaults

- Each container connected to a private virtual network “bridge”
- Each virtual network routes through NAT firewall on host IP
- All containers on virtual network can talk to each other without -p
- Best practice: to create a new virtual network for each app:
  - Network ‘my\_web\_app” for mysql and php/apache containers
  - Network “my\_api” for mongo and nodejs containers

# Docker Networks: CLI Management

- To show networks
  - ***docker network ls***
- Inspect a network
  - ***docker network inspect***
- Create a network
  - ***docker network create --driver***
- Attach a network to container
  - ***docker network connect***
- Detach a network from a container
  - ***docker network disconnect***

# Docker Networks: Default Security

- Create apps so frontend/backend sit on same Docker network
- Their intercommunication never leaves host
- All externally exposed ports closed by default
- You must manually expose via -p, which is better default security!



# Docker Image

# Section Overview:

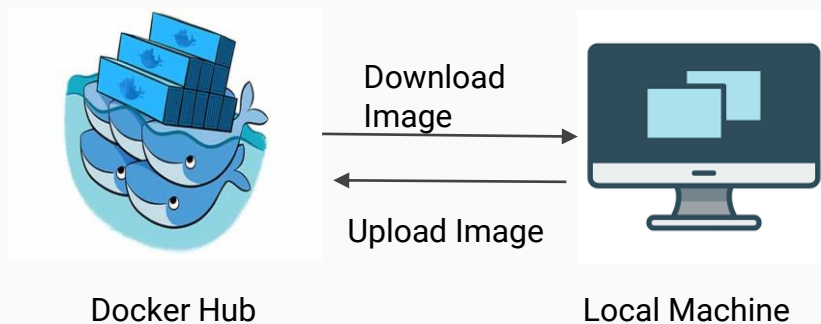
- What is an Image and image layers
- Docker Hub Registry: How to use it?
- Managing local image cache
- Build your own images

# What is a Docker Image?

- Images are made up of App binaries, dependencies and metadata
- Read-only template of instructions on creating a container
- "An Image is an ordered collection of root filesystem changes and the corresponding execution parameters for use within a container runtime."
- Each command is a layer in the image
- Does not contain a full OS
- Store in the Docker Engine Image cache
- Permanent Storage should be in Image registry
- Create your own images or use created by others and published in registry

# DockerHub

- Register at [hub.docker.com](https://hub.docker.com)
- When we create images on Docker Hub, we have to create them with our account name in front of them
- The images that have just the name of the repo are considered official.
  - Official Images: well tested, have quality documentation and obey best practice rules
- Images are tagged
  - A version of image can have more than one tag



# How to create Docker Image?

- Use **Dockerfile**
- Define steps to create an image and run it
- Each instruction in Dockerfile creates a layer in image
- Move images in/out of cache via:
  - Local filesystem via tarballs
  - push/pull to remote 'image registry' (e.g. Docker Hub)
- Images are not ideal for persistent data
  - Mount a host file system path into container
  - Use **docker volume** to create storage for persistent/unique data

# What is Dockerfile?



- A shell script
- Defines the commands/actions to be executed
- Example: Issuing ***docker run*** command starts the nginx server



→  
build



***docker build -t test***

- -t used to specify tag name for the image

# The Dockerfile Keywords:

- FROM (base image)
  - Every Dockerfile starts with FROM keyword
  - Specify the image name which will act as base image for the container
  - Example: *FROM nginx*
- ENV (environment variable)
  - Specify environment variables for the Docker Image
  - Example: *JAVA\_HOME*
- RUN (any arbitrary shell command)
  - Will be executed when we build the image

# The Dockerfile Keywords:

- **EXPOSE :**
  - open port from container to the virtual network
- **VOLUME:**
  - Define the storage location for the unique data generated by the application
- **CMD :**
  - command to run when the container starts
  - This will be executed when ***docker run*** is issued



# Image and its layers:

- Every image is a collection of layers.
- Each layer is uniquely identified and only stored once on a host
- This saves storage space on host and transfer time on push/pull
- container is just a single read/write layer on top of image
- ***docker image history*** and ***inspect*** commands can help us to get sense of image

## Assignment 2: Build Your Own Image

- Take existing Node.js app and dockerize it
- Make Dockerfile. Build it. Test it. Push it. Remove it. Run it.
- Details in ***part3-dockerfile-assignment/Dockerfile***
- Use the Alpine version of the official 'node' 6.x image
- Expected results is web site ***“https://localhost”***
- Tag and push to your Docker Hub account (free)
- Remove your image from local cache, run again from Hub

# Docker Volumes

# Section Overview:

- Understanding the Persistent Data Problem
- Learn and attach volumes to the container
- Learn about Bind Mounting

# Defining problem of persistent data:

- Containers are usually immutable and ephemeral
  - Immutable: a container won't be modified during its life: no updates, no patches, no configuration changes
  - Containers can be restarted, stopped or replaced easily
  - Redeploy a new container, if configuration needs to change
- Persistent Data problem: What about the unique data produced by application?
- Two solutions:
  - Volumes and Bind Mounts

# Volumes: Making Docker persistent

- Volumes: Make special location outside of container UFS
- **VOLUME** command in **Dockerfile**
- Also override with ***docker run -v /path/in/container***
- Bypasses Union File System and stores in alt location on host
- Includes it's own management commands under ***docker volume***
- Connect to none, one or multiple containers at once
- Not subject to ***commit***, ***save***, or ***export*** commands
- By default they only have a unique ID, but you can make it “named volume” by assigning name to it.

# Bind Mounting

- Bind Mounts: link container path to host path
- Maps a host file or directory to a container file or directory
- Basically just two locations pointing to the same file(s)
- Again, skips UFS, and host files overwrite any in container
- Can't use in Dockerfile, must be at *container run*
- *... run -v /Users/bret/stuff:/path/container (mac/linux)*
- *... run -v //c/Users/bret/stuff:/path/container(windows)*

# Docker Compose



# What is Docker Compose?

- A tool for defining and running multi-container docker applications
  - A YAML file is used to configure application services
  - With single command, create and start all the services from configuration
- Comprised of 2 separate but related things:
  - YAML-formatted file that describes our solution options for: **containers**, **networks** and **volumes**
  - A CLI tool docker-compose used for local dev/test automation with those YAML files
- For more learning on docker compose:
  - <https://docs.docker.com/compose/>

# docker-compose.yml

- Compose YAML format has its own versions: 1, 2, 2.1, 3, 3.1
- YAML file can be used with *docker-compose* command for local docker automation
- `docker-compose --help`
- *docker-compose.yml* is default filename, but any can be used with *docker-compose -f*

# docker-compose CLI

- CLI tool comes with Docker for Windows/Mac, but separate download for Linux
- Not a production-grade tool but ideal for local development and test
- Two most common commands are
  - *docker-compose up* # setup volumes/networks and start all containers
  - *docker-compose down* # stop all containers and remove cont/vol/net
- If all your projects had a Dockerfile and docker-compose.yml then "new developer onboarding" would be:
  - *git clone github.com/some/software*
  - *docker-compose up*

## Assignment 3: Writing A Compose File

- Build a basic compose file for a Drupal content management system website.
  - Use Docker Hub
- Use the *drupal* image along with the *postgres* image
- Use *ports* to expose **Drupal on 8080** so you can **localhost:8080**
- Be sure to set **POSTGRES\_PASSWORD** for postgres
- Walk through Drupal setup via browser
- Tip: Drupal assumes DB is *localhost*, but it's service name
- **Extra Credit:** Use volumes to store Drupal unique data

# Additional Resources:

- [Docker Homepage](#)
- [Docker Hub](#)
- [Docker Blog](#)
- [Docker Documentation](#)
- [Docker Getting Started Guide](#)
- [Docker Help](#) on StackOverflow